Insights to HDInsight

Ioannis Stavrinides, Microsoft

Why Hadoop in the Cloud?

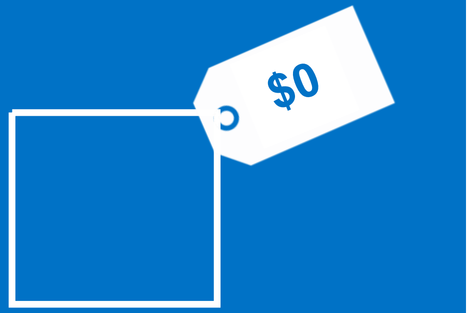

No HW costs

Unlimited scale

Pay what
you need

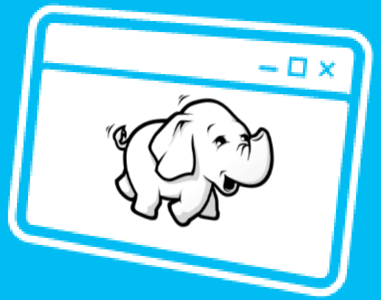

Deployed
in minutes

No hardware costs

Hadoop in the Cloud bypasses hardware costs

Hardware acquisition

Hardware maintenance

Performance tuning

No HW costs

Unlimited Scale

Hadoop in the Cloud bypasses capacity planning

Spin up any number of Hadoop nodes on-demand

Go from tens of nodes to thousands of nodes

Pay for What You Need

Hadoop is billed by usage

Billed for usage

Clusters can be deleted when no longer used

Pay what
you need

Deployed in minutes

- Hadoop in the Cloud Bypasses deployment expertise
  - Hadoop is non-trivial to install and get up and running on multi-nodes
  - Education gap in IT community regarding Hadoop
- Hadoop is deployed in minutes
  - Spin up any number of Hadoop nodes on-demand
  - Up and running in a few clicks (and within minutes)

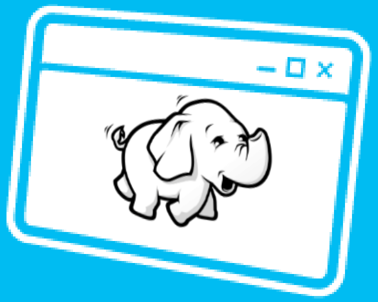

Deployed
in minutes

Azure HDInsight

Big Data made easy

Enterprise Ready

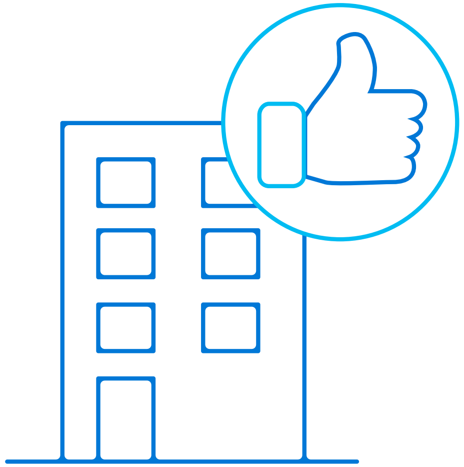

Easier and more productive for all users

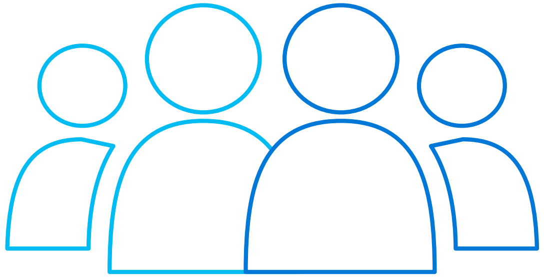

Hybrid

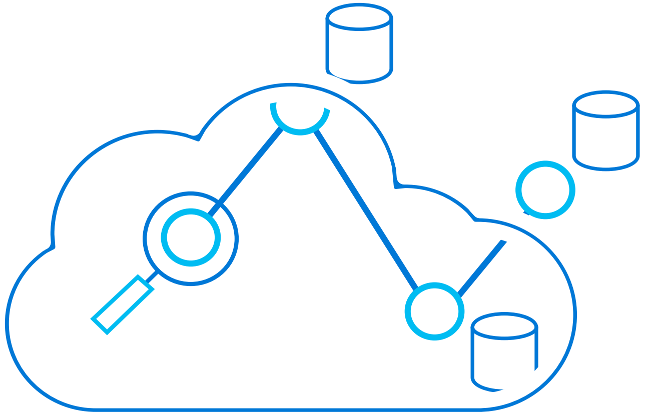

Azure HDInsight

Big Data made easy

Enterprise Ready

Easier and more productive for all users

Hybrid

Highly Available – Designed for the cloud ground up

- HDInsight provides primary and secondary headnodes allowing for better reliability
- Have invested in making entire stack including Resource Manager, HiverServer2 HA ready
- HDInsight stack includes Zookeeper nodes at no extra charge to customer

Highest availability guarantee in the industry for peace of mind

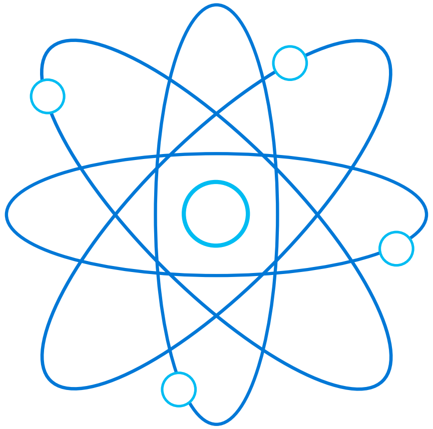

99.9% SLA

- Managed, monitored and supported by Microsoft
- Enterprise-leading SLA—99.9% uptime for both VM connectivity and Hadoop running in VMs
- No IT resources needed for upgrades and patching
- Microsoft monitors your deployment so you don’t have to

*Applies to HDInsight only

Always encrypted, Role-based security & Auditing

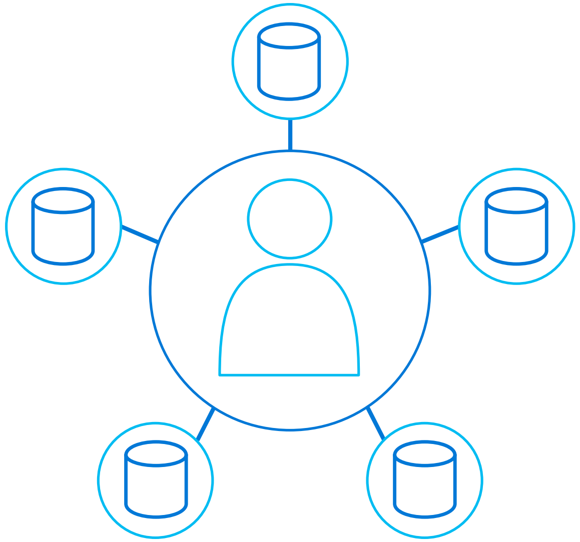

- Always encrypted; in motion using SSL, and at rest using keys in Azure Key Vault
- Single sign-on, multi-factor authentication and integration of on-premises identities w/Active Directory integration
- Fine-grained ACLs for role-based access controls with Apache Ranger
- Auditing every access / configuration change with Apache Ranger

Alerting, monitoring, and pre-emptive actions

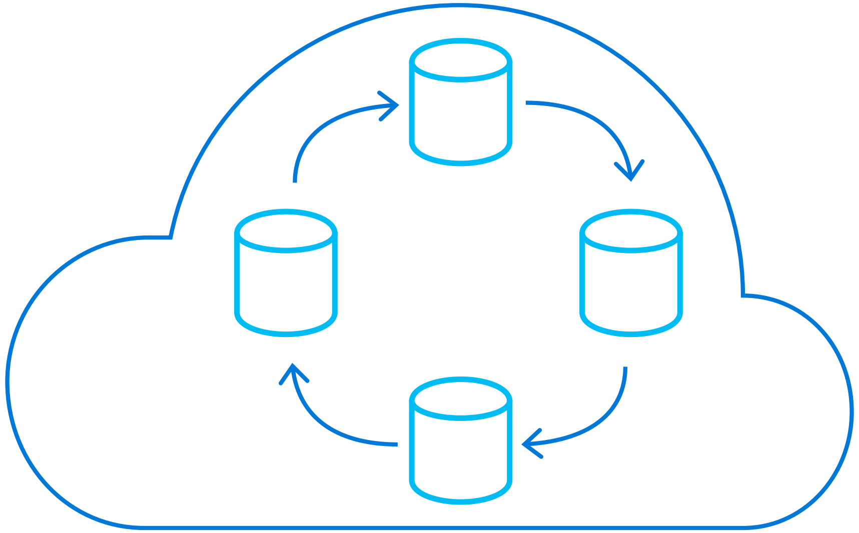

- Enhanced workload protection through integration with Microsoft Operations Management Suite (OMS)
- Threat detection, monitoring, and management

Petabyte size files and Trillions of objects

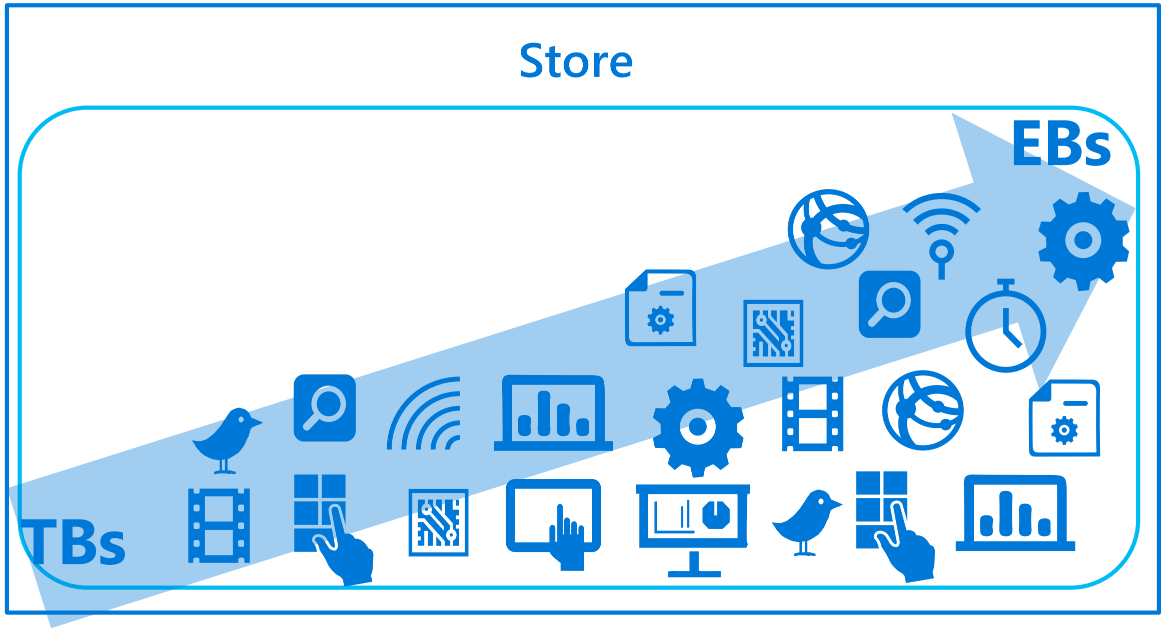

- Store data in it’s native format
- PB sized files, 200x larger than anyone else
- Scalable throughput for massively parallel analytics
- No need to redesign application or reparation data at higher scale

Backed by Microsoft and Hortonworks

Microsoft

+

HORTONWORKS®
POWERING THE FUTURE OF DATA™

- Microsoft + Hortonworks has **37 committers** for Hadoop Core; more than all managed cloud Hadoop vendors combined
- Uniquely ready to support your deployment
- Can fix and commit code back to Hadoop

Runs in the most datacenters worldwide

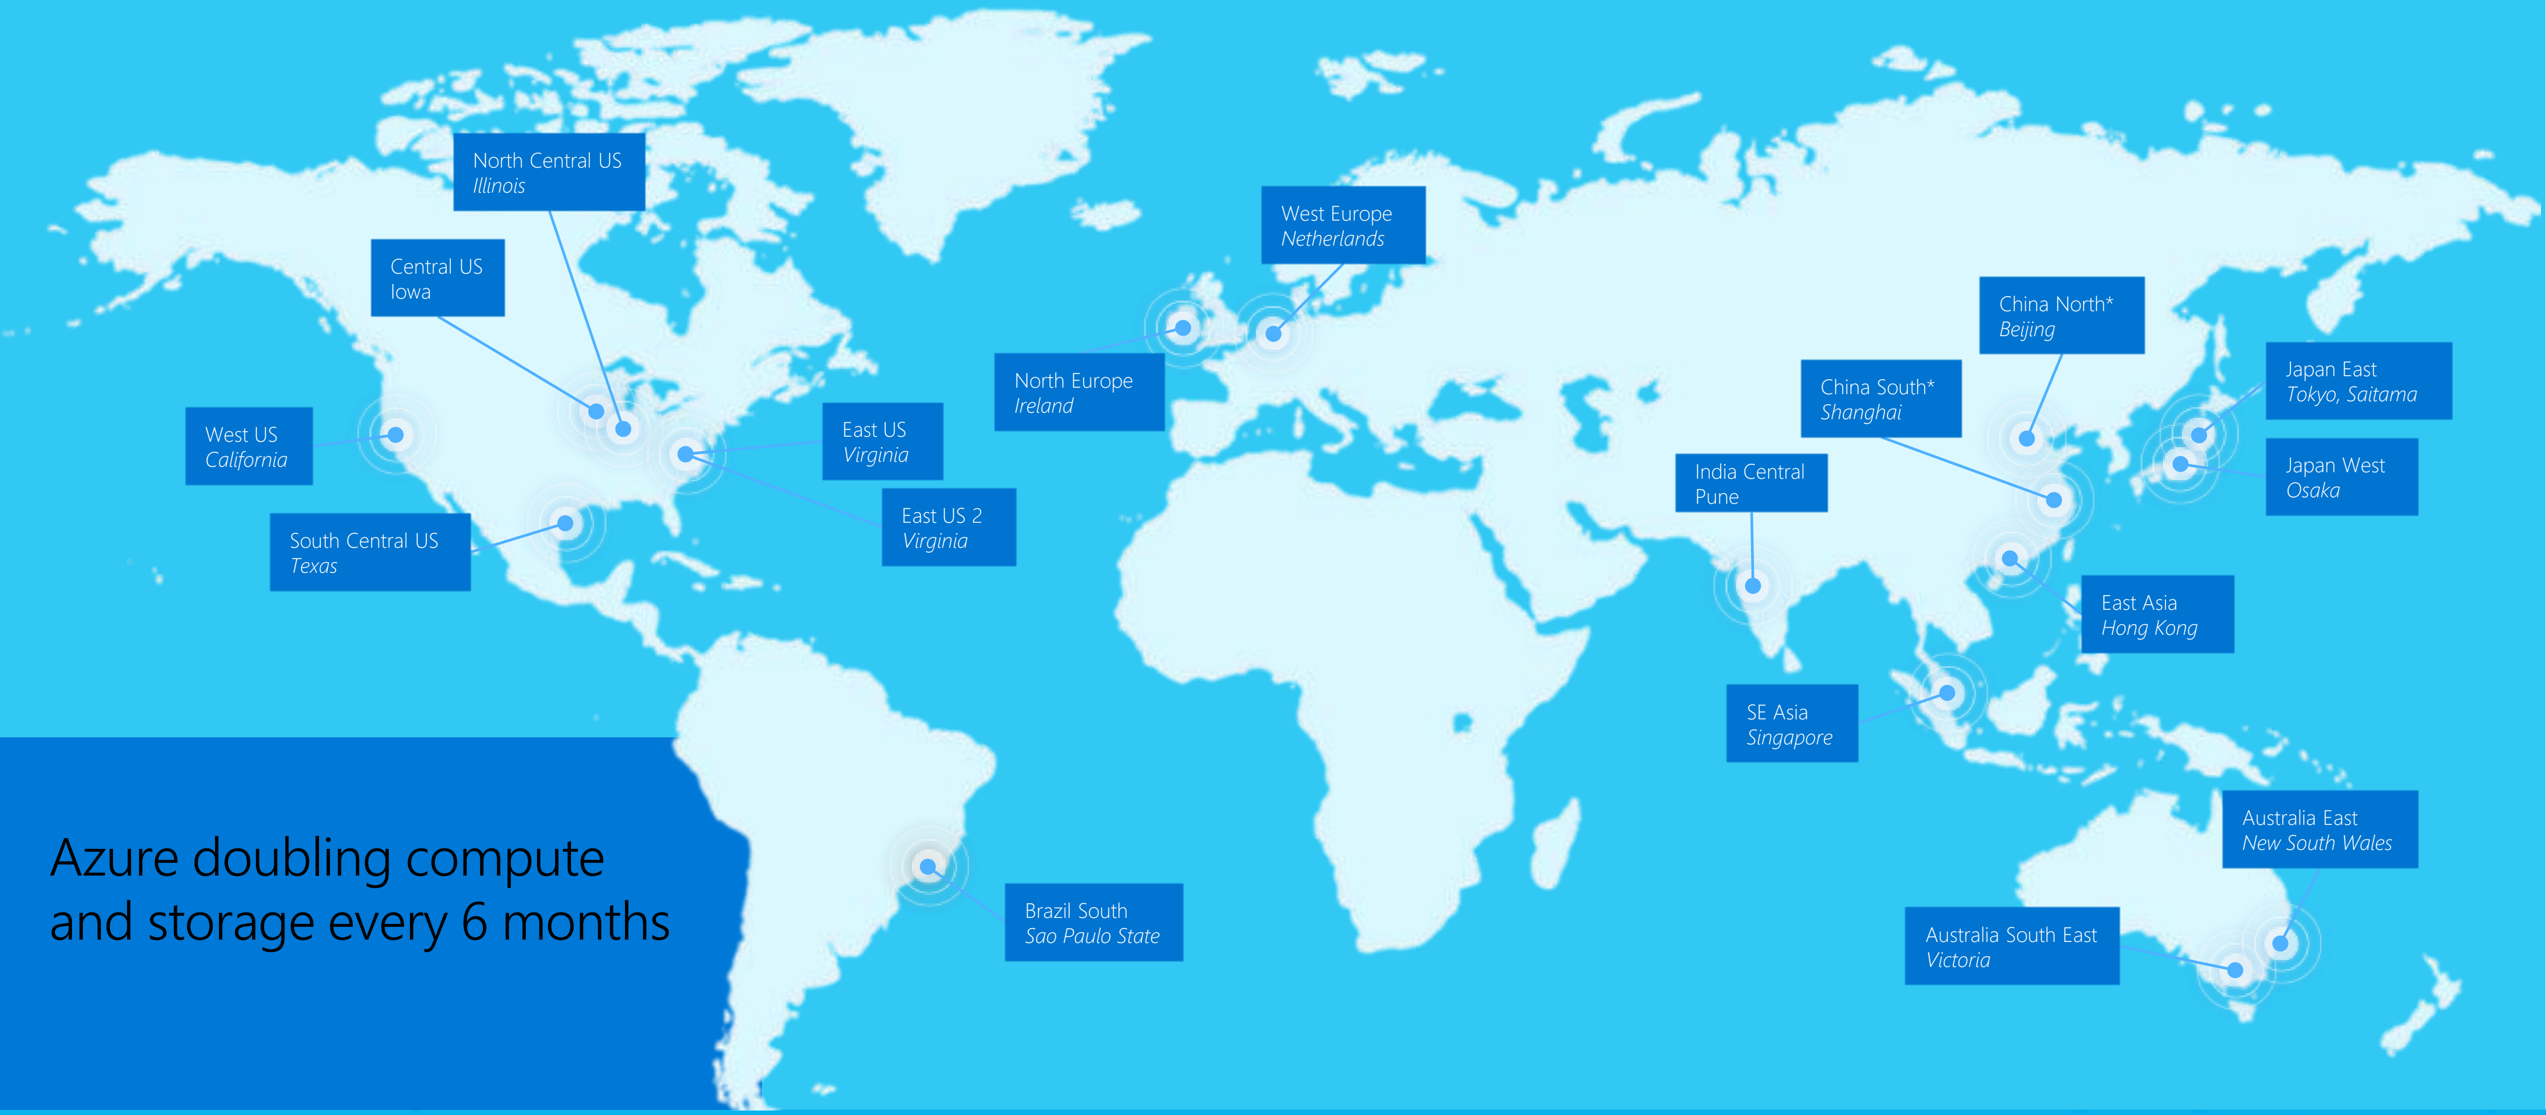

North Central US
Illinois

Central US
Iowa

West US
California

South Central US
Texas

East US
Virginia

East US 2
Virginia

North Europe
Ireland

West Europe
Netherlands

Brazil South
Sao Paulo State

India Central
Pune

China South*
Shanghai

China North*
Beijing

Japan East
Tokyo, Saitama

Japan West
Osaka

East Asia
Hong Kong

SE Asia
Singapore

Australia East
New South Wales

Australia South East
Victoria

Azure doubling compute and storage every 6 months

Lower total cost of ownership

- No hardware
- Hadoop support included with Azure support
- Pay only for what you use
- Independently scale storage and compute
- No need to hire specialized operations team
- 63% lower total cost of ownership than on-premises*

*IDC study "The Business Value and TCO Advantage of Apache Hadoop in the Cloud with Microsoft Azure HDInsight"

Azure HDInsight

Big Data made easy

Enterprise Ready

Easier and more productive for all users

Hybrid

Easy for administrators to spin up quickly

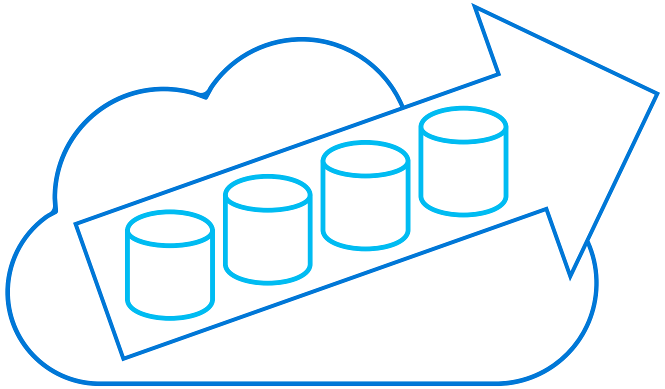

- Deploy big data projects in minutes
- No hardware to install, tune, configure or deploy
- No infrastructure or software to manage
- Scale to tens to thousands of machines instantly

Debug and Optimize your Big Data programs with ease

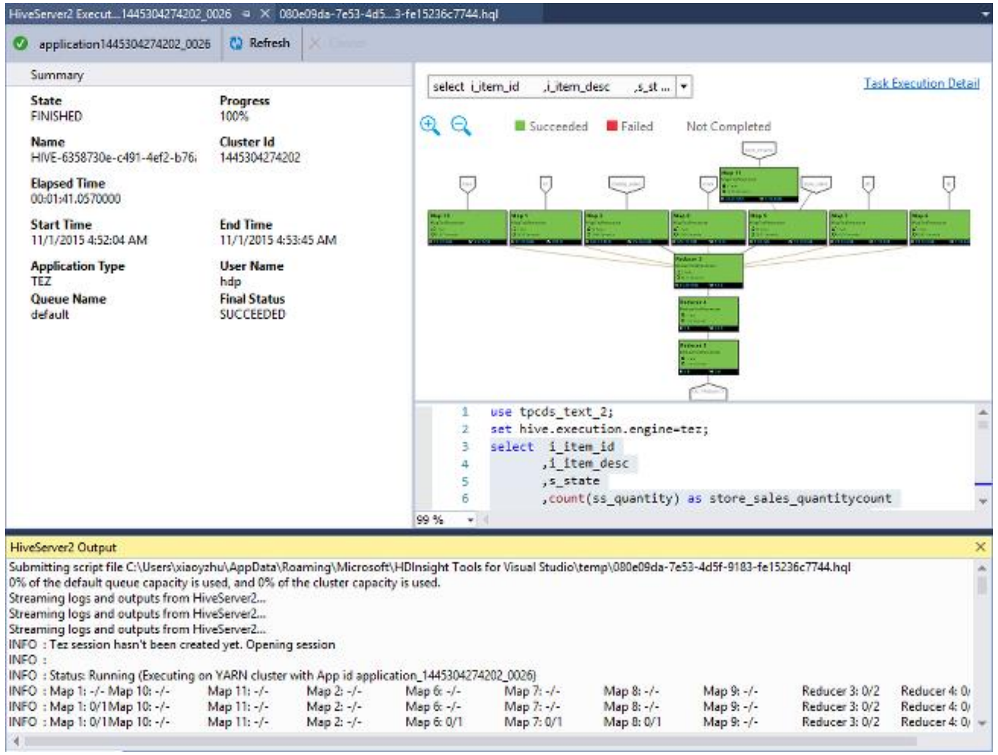
HiveServer2 Execut...1445304274202_0026
080e09da-7e53-4d5...3-fe15236c7744.hql
application1445304274202_0026
Refresh
Cancel
Summary
select i_item_id ,i_item_desc ,s_st ...
Task Execution Detail
🟩 Succeeded
🟥 Failed
Not Completed
99 %

State
FINISHED

Progress
100%

Name
HIVE-6358730e-c491-4ef2-b76

Cluster Id
1445304274202

Elapsed Time
00:01:41.0570000

Start Time
11/1/2015 4:52:04 AM

End Time
11/1/2015 4:53:45 AM

Application Type
TEZ

User Name
hdp

Queue Name
default

Final Status
SUCCEEDED

```
use tpcds_text_2;
set hive.execution.engine=tez;
select  i_item_id
       ,i_item_desc
       ,s_state
       ,count(ss_quantity) as store_sales_quantitycount
```

HiveServer2 Output

```
Submitting script file C:\Users\xiaoyzhu\AppData\Roaming\Microsoft\HDInsight Tools for Visual Studio\temp\080e09da-7e53-4d5f-9183-fe15236c7744.hql
0% of the default queue capacity is used, and 0% of the cluster capacity is used.
Streaming logs and outputs from HiveServer2...
Streaming logs and outputs from HiveServer2...
Streaming logs and outputs from HiveServer2...
INFO  : Tez session hasn't been created yet. Opening session
INFO  :
INFO  : Status: Running (Executing on YARN cluster with App id application_1445304274202_0026)
INFO  : Map 1: -/- Map 10: -/-    Map 11: -/-    Map 2: -/-    Map 6: -/-    Map 7: -/-    Map 8: -/-    Map 9: -/-    Reducer 3: 0/2    Reducer 4: 0/
INFO  : Map 1: 0/1Map 10: -/-     Map 11: -/-    Map 2: -/-    Map 6: -/-    Map 7: -/-    Map 8: -/-    Map 9: -/-    Reducer 3: 0/2    Reducer 4: 0/
INFO  : Map 1: 0/1Map 10: -/-     Map 11: -/-    Map 2: -/-    Map 6: 0/1    Map 7: 0/1    Map 8: 0/1    Map 9: -/-    Reducer 3: 0/2    Reducer 4: 0/
```

- Deep integration with IDEs for developer productivity: Visual Studio, Eclipse, & IntelliJ
- Integrated with Hive, Pig, Storm, and Spark
- Visually see execution of Hive jobs ran by the Tez execution engine
- Full Intellisense

Easy notebook experience for data engineers

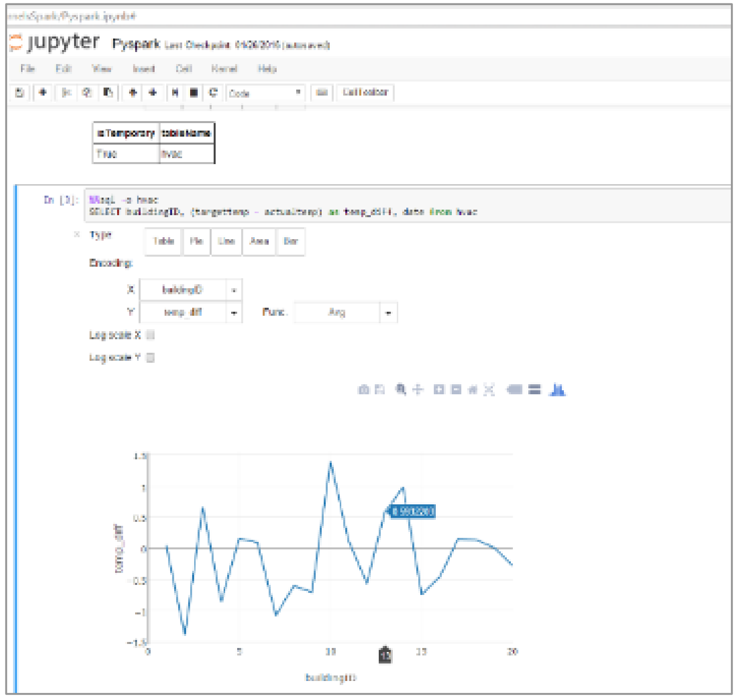

isTemporary	tableName
True	hvac

```
In [3]: %%sql -o hvac
        SELECT buildingID, (targettemp - actualtemp) as temp_diff, date from hvac
```

- Most popular notebooks, Jupyter and Zeppelin out-of-the-box
- Combine code, statistical equations and visualizations
- Worked w/ Jupyter community to enhance kernel to allow Spark execution through REST endpoint

Easy for data scientists with familiar R language

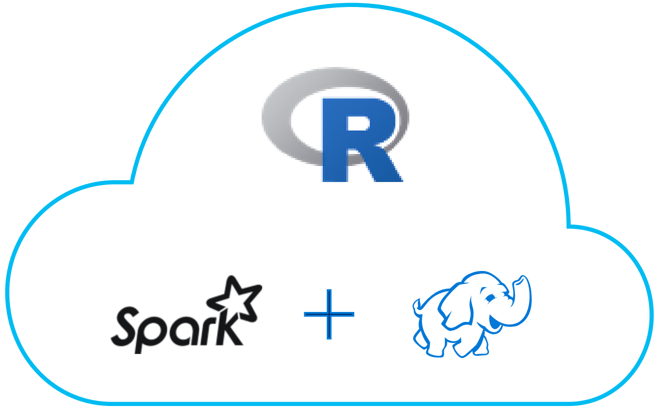

*Applies to HDInsight only

R Server for HDInsight

- Largest portable R parallel analytics library
- Terabyte-scale machine learning—1,000x larger than in open source R
- Up to 100x faster performance using Spark and optimized vector/math libraries
- Enterprise-grade security and support

Easy for business analysts with interactive reports over big data

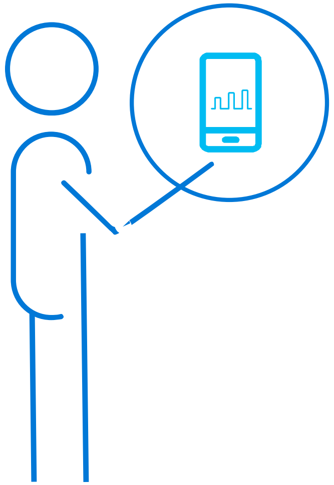

- Interactive BI with big data
- Spark 2.0 integration
- Interactive Hive with LLAP- keeps data compressed running in-memory 25x faster
- ODBC driver to use Power BI or third party tools (Tableau, SAP, Qlik, etc.)

Azure HDInsight

Big Data made easy

Enterprise Ready	Easier and more productive for all users	Hybrid

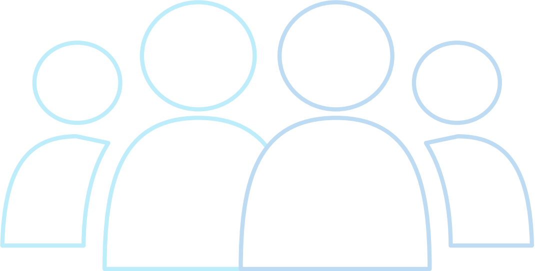

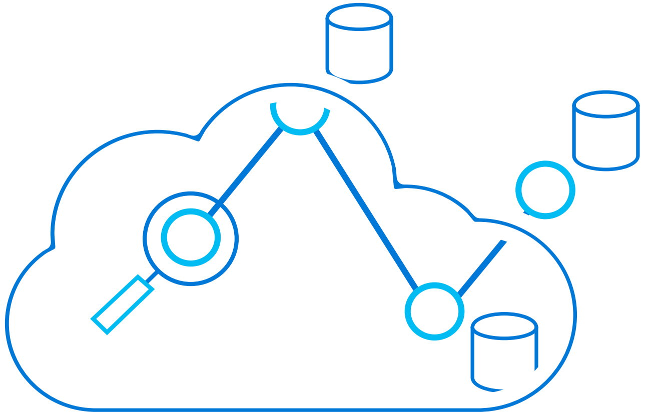

On-premises and cloud

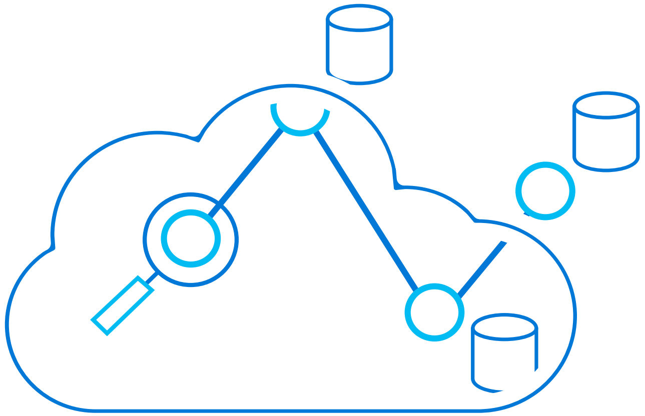

- Uses Hortonworks Data Platform (HDP)
- Move projects from on-premises to cloud without code rewrite
- Hybrid scenarios supported like Dev/Test, burst, back up, disaster recovery

Recognized by top analysts

FORRESTER®

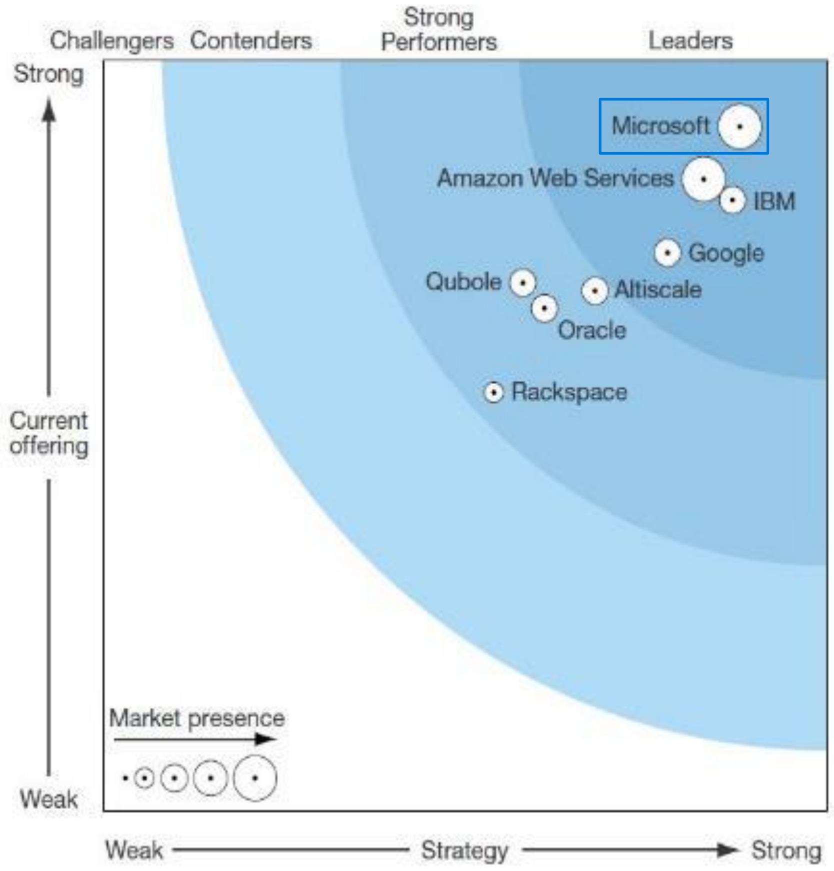

Forrester Wave for Big Data Hadoop Cloud

- Named industry leader by Forrester with the most comprehensive, scalable, and integrated platforms*
- Recognized for its cloud-first strategy that is paying off*

*The Forrester WaveTM: Big Data Hadoop Cloud Solutions, Q2 2016.

Hadoop Workloads

Microsoft

Hadoop is a platform with portfolio of projects

- Governed by Apache Software Foundation (ASF)
- Comprises core services of MapReduce, HDFS, and YARN
- In addition to the core, includes functions across:
  - Governance and integration, Tools, Data Access, Security, and Operations

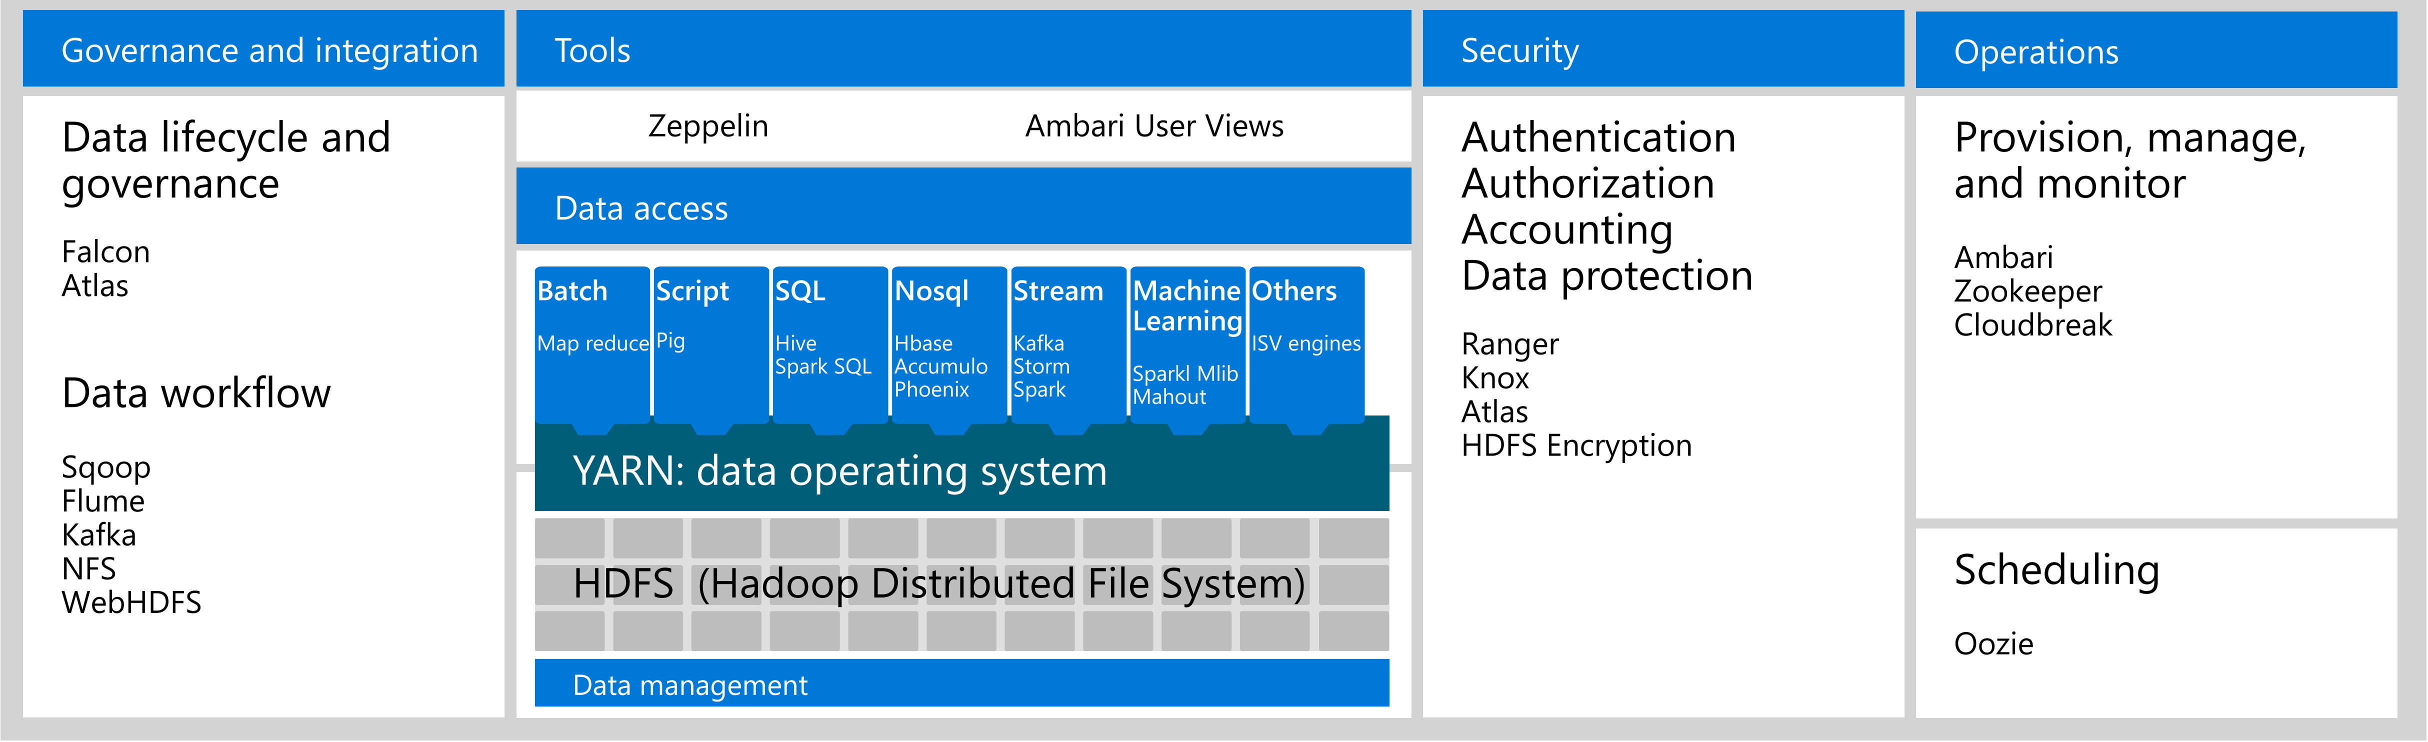

Governance and integration

Data lifecycle and governance

Falcon
Atlas

Data workflow

Sqoop
Flume
Kafka
NFS
WebHDFS

Tools

Zeppelin

Ambari User Views

Data access

Batch	Script	SQL	Nosql	Stream	Machine Learning	Others
Map reduce	Pig	Hive Spark SQL	Hbase Accumulo Phoenix	Kafka Storm Spark	Sparkl Mlib Mahout	ISV engines

YARN: data operating system

HDFS (Hadoop Distributed File System)

Data management

Security

Authentication
Authorization
Accounting
Data protection

Ranger
Knox
Atlas
HDFS Encryption

Operations

Provision, manage, and monitor

Ambari
Zookeeper
Cloudbreak

Scheduling

Oozie

HDFS

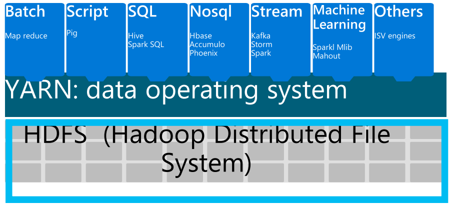

Batch

Map reduce

Script

Pig

SQL

Hive
Spark SQL

Nosql

Hbase
Accumulo
Phoenix

Stream

Kafka
Storm
Spark

Machine Learning

Sparkl Mlib
Mahout

Others

ISV engines

YARN: data operating system

HDFS (Hadoop Distributed File System)

HDFS is a distributed file system

From a few nodes to thousands of nodes

Files can be spread out over multiple nodes

HDFS stores large amounts of data

Very large files are supported including those larger than the capacity of a single node

HDFS stores non-relational files

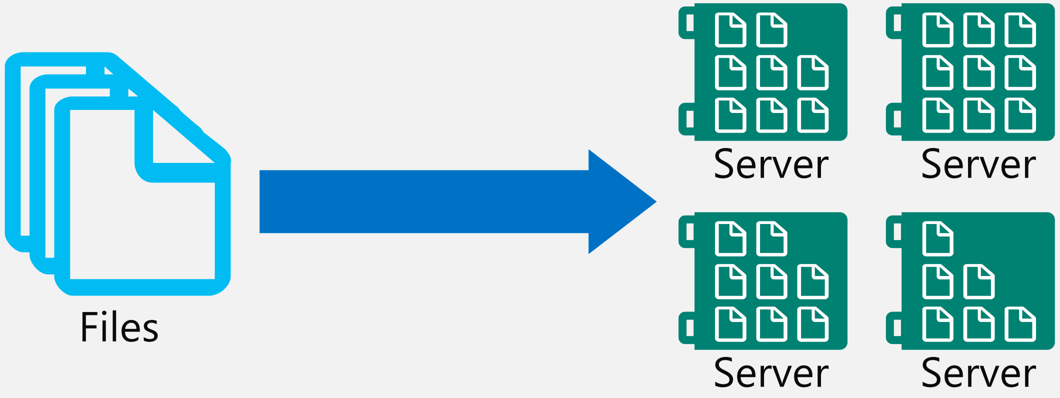

Files

Server

Server

Server

Server

MapReduce

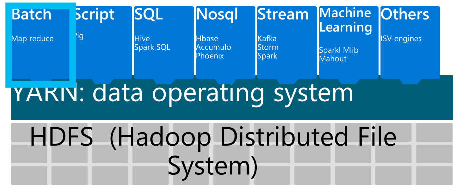

Takes processing to where data is

Distributed processing: instead of serializing processing through one pipe, distributes computing locally where data is

Brings back only the resultant data

Scales linearly as you add nodes

Three-step execution

Map: Developer writes map functions to the data

Shuffle / Distributes: Framework automatically shuffles for you (networking, synchronization, recovery, scheduling)

Reduce: Developer writes reduce functions to bring resultant data back

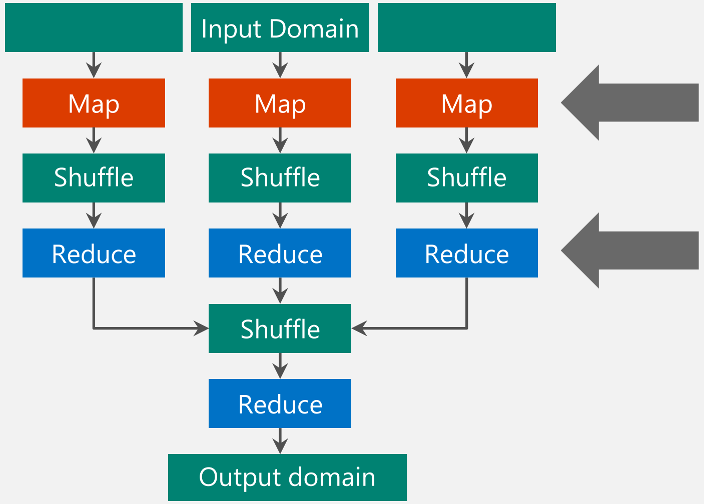

```
// Map Reduce function in JavaScript

var map = function (key, value, context) {
var words = value.split(/[^a-zA-Z]/);
for (var i = 0; i < words.length; i++) {
                if (words[i] !== "")
context.write(words[i].toLowerCase(),
1);}
}};

var reduce = function (key, values, context) {
var sum = 0;
while (values.hasNext()) {
sum += parseInt(values.next());
        }
context.write(key, sum);
};
```

Hive

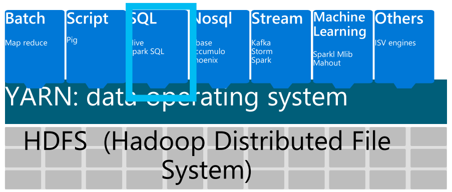

SQL-like queries on Hadoop data in HDFS

HiveQL is a SQL-like language (subset of SQL)

Hive structures include well-understood database concepts such as tables, rows, columns, partitions

Compiled into MapReduce jobs that are executed on Hadoop

Dramatic performance gains with Hive w/LLAP

Performance gains up to 25x

ODBC drivers to integrate with Power BI, Tableau, Qlik, etc.

Opens up scenarios to do interactive BI and reporting on big data

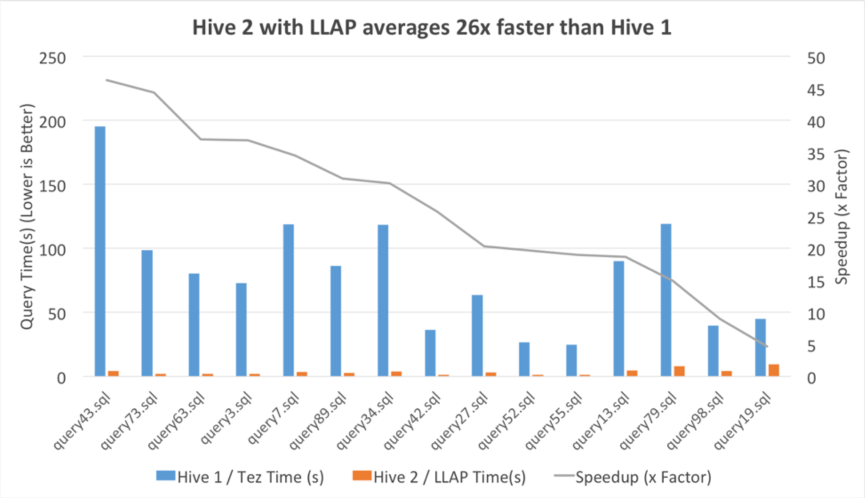

Hive 2 with LLAP averages 26x faster than Hive 1

HBase

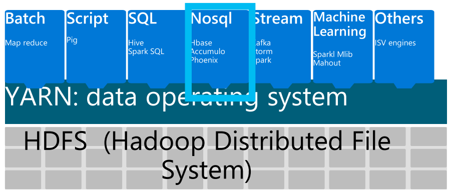

NoSQL database on data in HDFS

Columnar, NoSQL database

Runs on top of the Hadoop Distributed File System (HDFS)

Provides flexibility in that new columns can be added to column families at any time

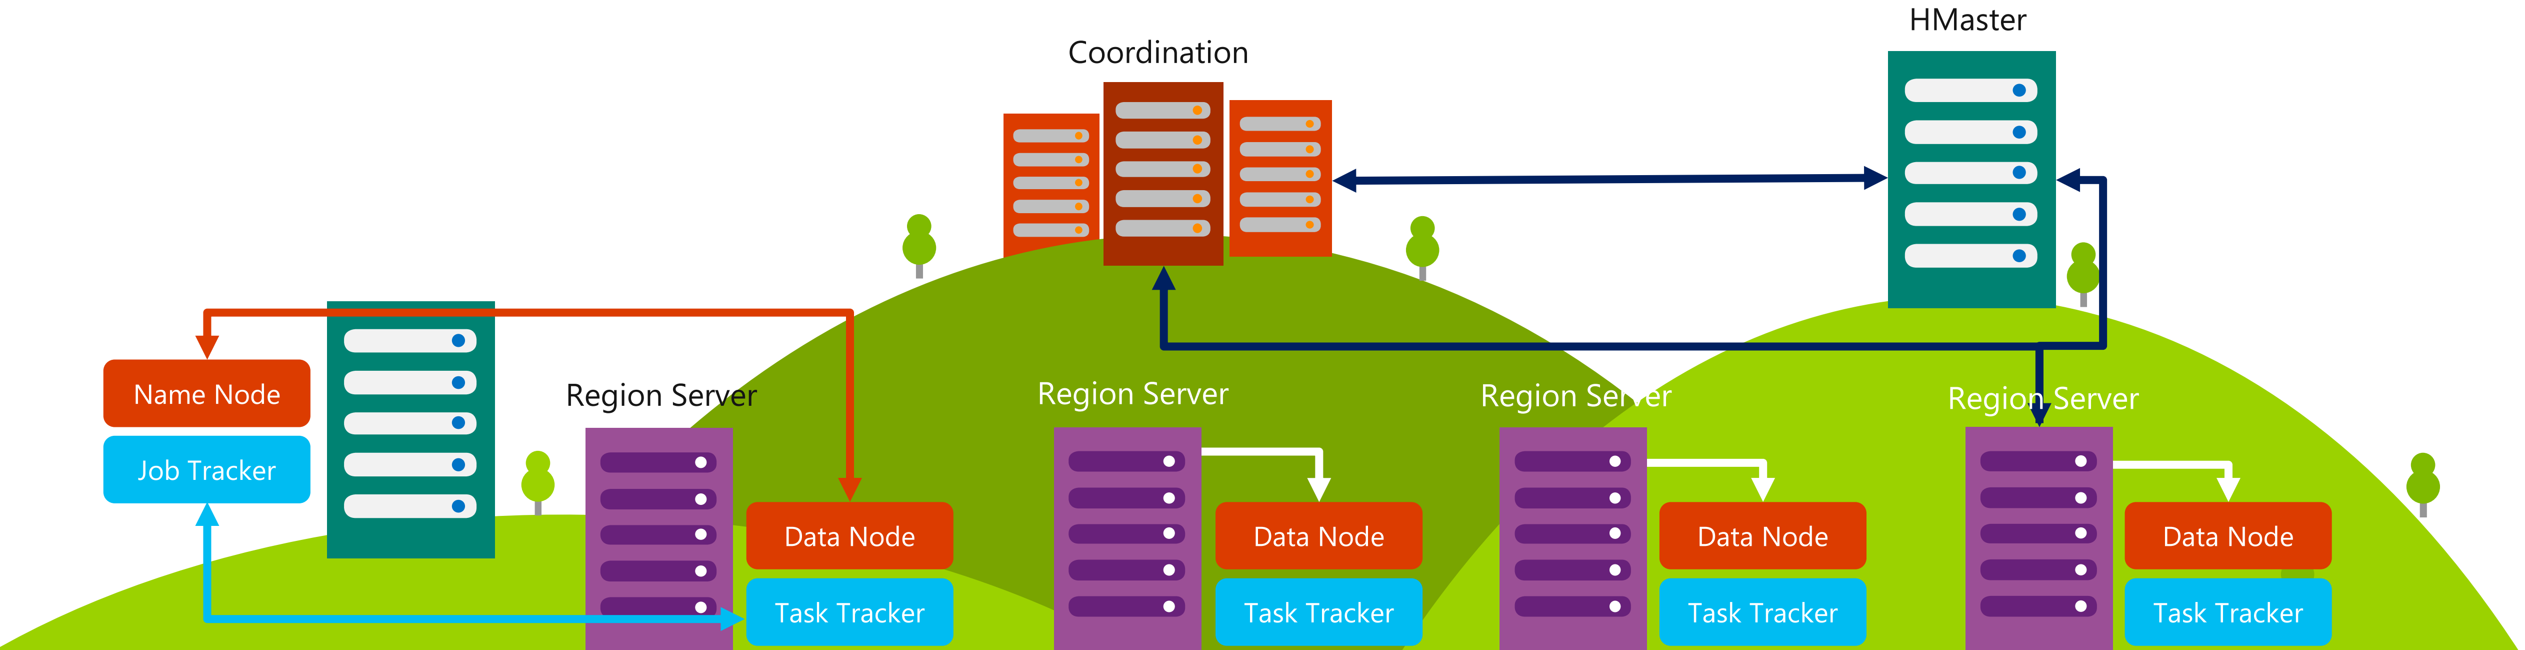

Storm

STORM

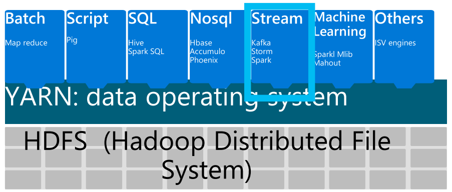

Batch
Map reduce

Script
Pig

SQL
Hive
Spark SQL

Nosql
Hbase
Accumulo
Phoenix

Stream
Kafka
Storm
Spark

Machine Learning
Sparkl Mlib
Mahout

Others
ISV engines

YARN: data operating system

HDFS (Hadoop Distributed File System)

Stream analytics for Near-Real Time processing

Consumes millions of real-time events from a scalable event broker (i.e.; Apache Kafka, Azure Event Hub)

Performs time-sensitive computation

Output to persistent stores, dashboards or devices

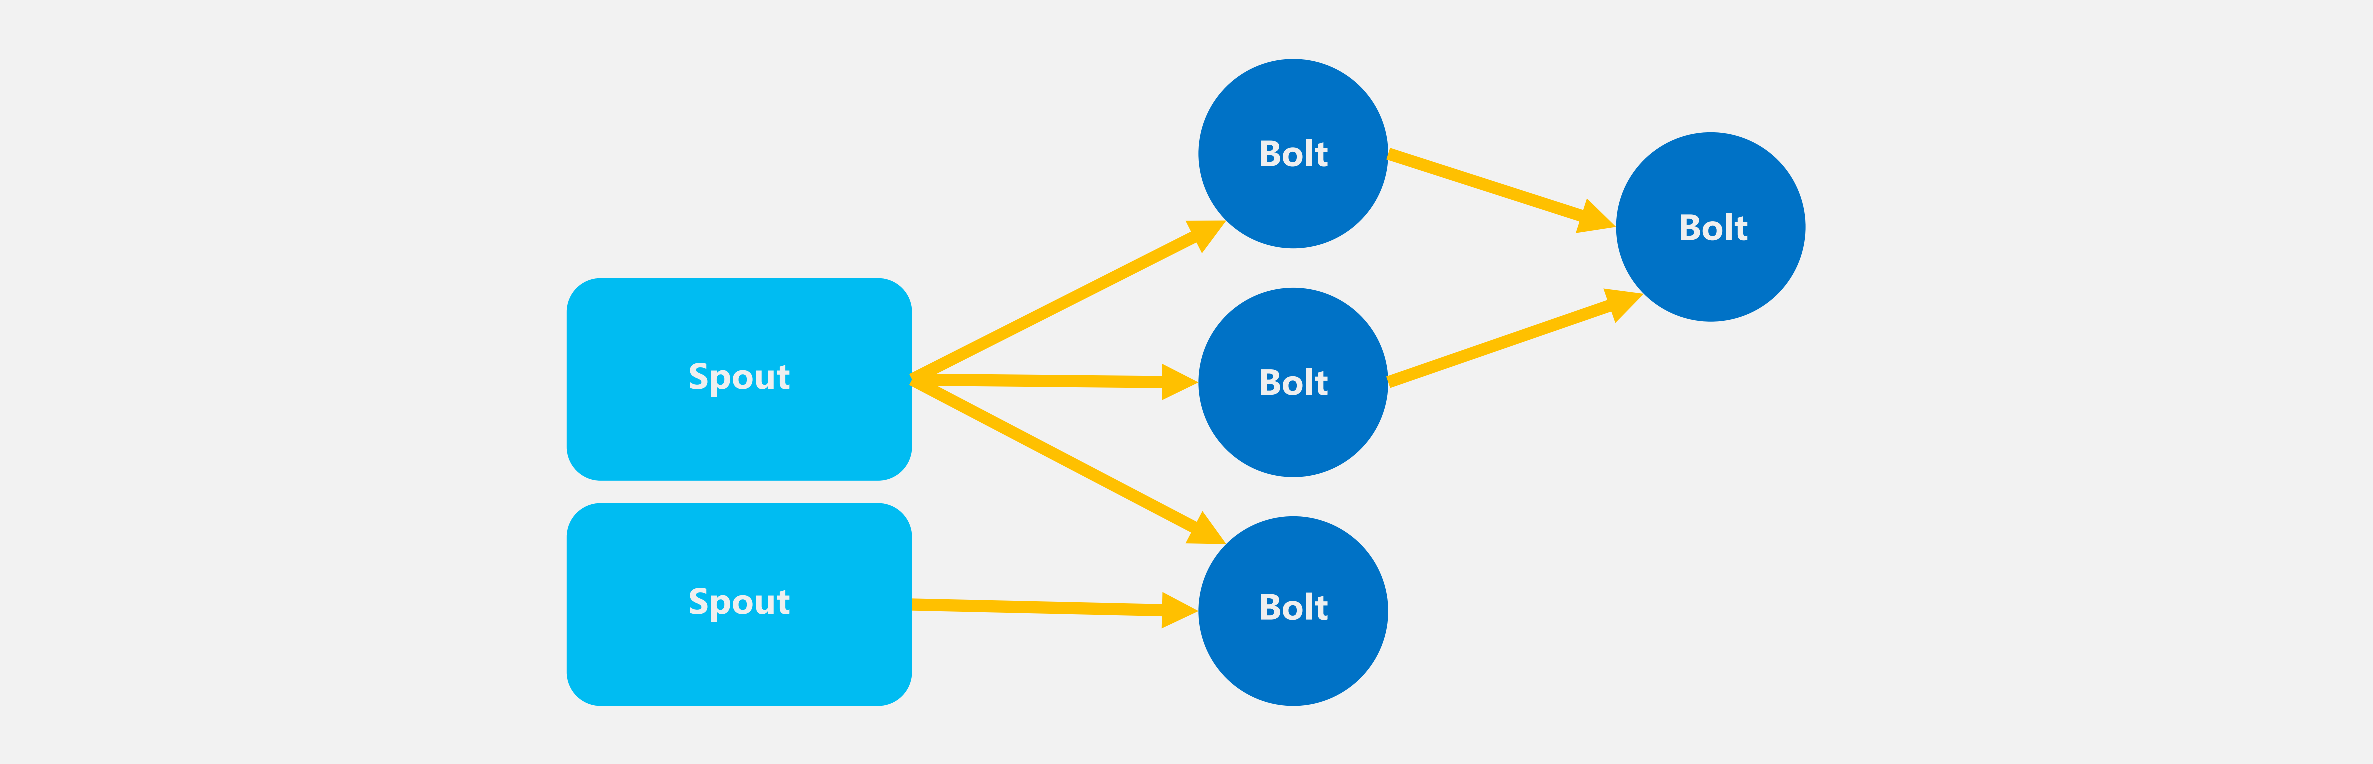

Spout

Spout

Bolt

Bolt

Bolt

Bolt

Kafka

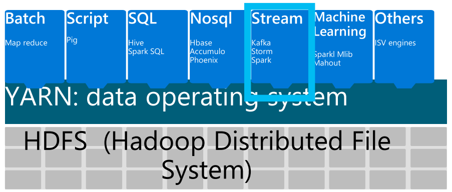

High-throughput, low-latency for real-time data

Stream millions of events per second

Enterprise-grade management and control

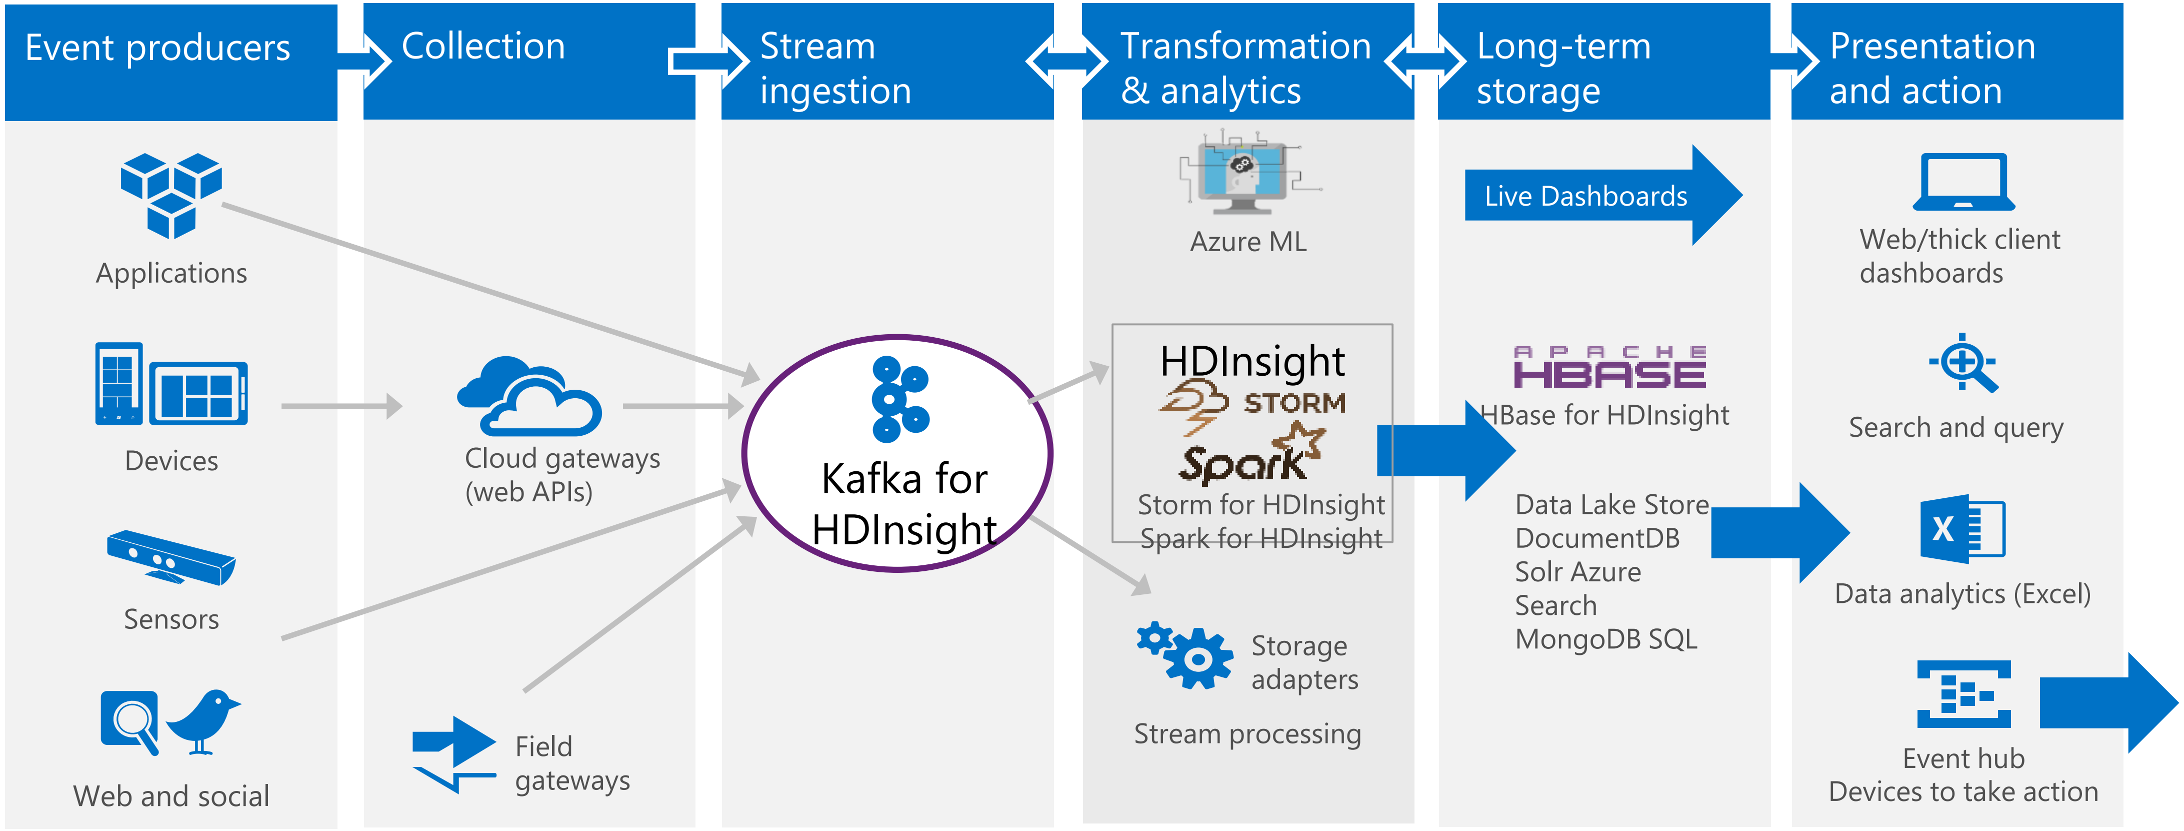

Mahout

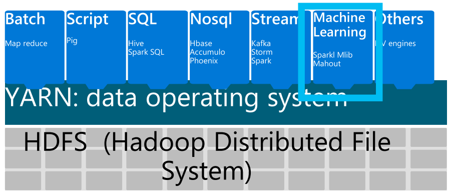

Machine learning library

A library of machine learning algorithms to execute on data in HDFS

Algorithms are not dependent on size of data and can scale with large datasets

Library includes: Collaborative Filtering, Classification, Clustering, Dimensionality Reduction, Topic Models

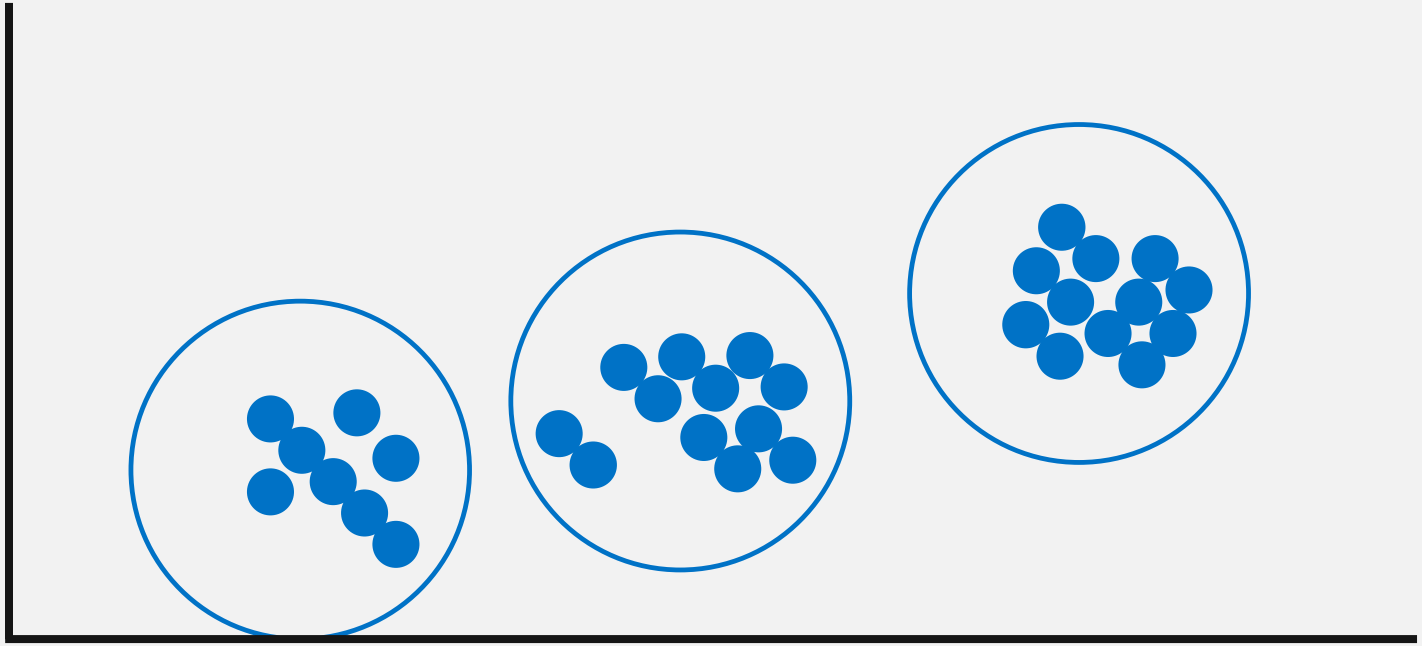

Spark

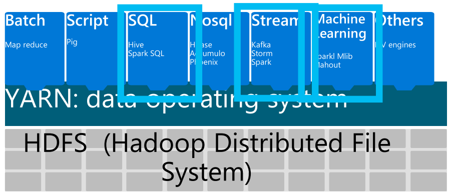

Massive data processing framework built on in-memory

Single execution model for multiple tasks

Processing up to 100x faster performance

Developer friendly (Java, Python, Scala)

BI tool of choice (Power BI, Tabelau, Qlik, SAP)

Notebook experience (Jupyter & Zeppelin)

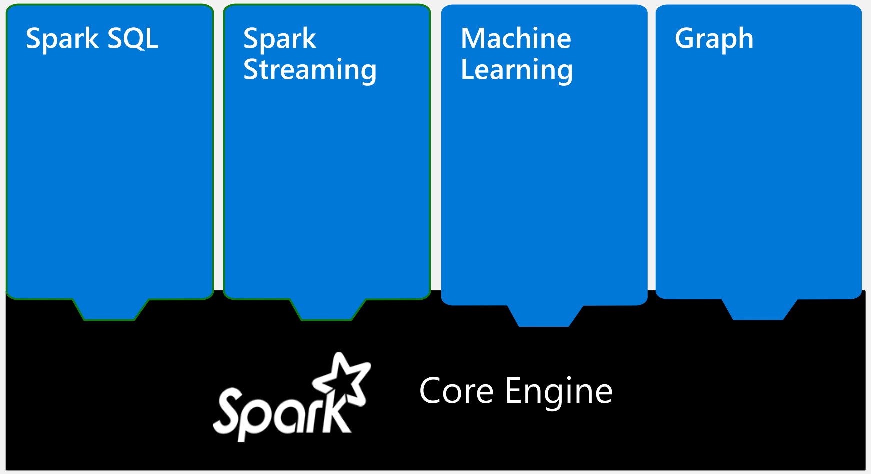

R Server

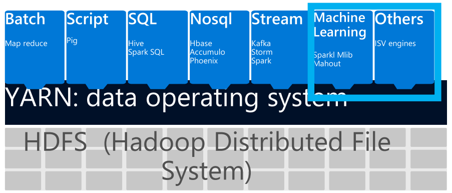

Batch
Map reduce

Script
Pig

SQL
Hive
Spark SQL

Nosql
Hbase
Accumulo
Phoenix

Stream
Kafka
Storm
Spark

Machine Learning
Sparkl Mlib
Mahout

Others
ISV engines

YARN: data operating system

HDFS (Hadoop Distributed File System)

Predictive analytics, machine learning, and statistical modeling for big data

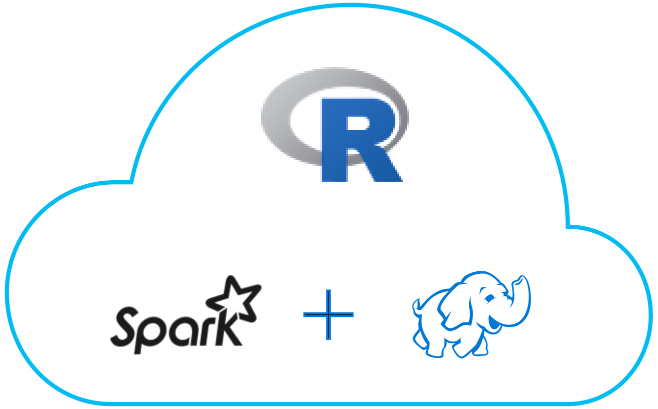

- Largest portable R parallel analytics library
- Terabyte-scale machine learning—1,000x larger than in open source R
- Up to 100x faster performance using Spark and optimized vector/math libraries
- Enterprise-grade security and support

ISV Integration

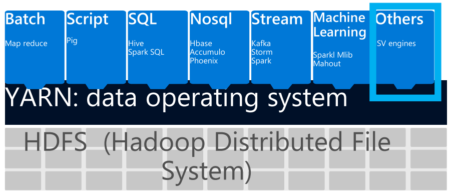

Integration with leading productivity applications

Spin up Hadoop and Spark clusters pre-integrated and pre-tuned with ISV applications out-of-the-box

Runs on the HDInsight clusters; does not require separate VMs

Fast and easy way to spin up applications

atscale

StreamSets™

Datameer®

CASK

Demo

Spark on HDInsight

Business in action with Cloud solutions

Digital Business Conference, Malta 2017